

VentureMiner AI

Product Requirements Document (PRD)

Document: 01-PRD / 01_product_requirements_document

Version: v1.0

Date: 2026-07-20

Author: VentureMiner AI Documentation Team

Status: Approved

Project: VentureMiner AI — AI Venture Intelligence Platform

Table of Contents

Product Requirements Document (PRD)

Table of Contents

1. Executive Summary
 - 1.1 One-line value proposition
 - 1.2 Strategic positioning
 - 1.3 At a glance
2. Product Vision & Mission
 - 2.1 Vision
 - 2.2 Mission
 - 2.3 Guiding principles
3. Market Opportunity
 - 3.1 Total Addressable Market (TAM)
 - 3.2 Serviceable Addressable Market (SAM)
 - 3.3 Serviceable Obtainable Market (SOM)
 - 3.4 Why now
 - 3.5 Competitive landscape
4. Problem Statement
 - 4.1 Jobs-to-be-done
 - 4.2 Today, this is broken
 - 4.3 Why this is the moment to fix it
5. Stakeholders & User Personas
 - 5.1 Primary personas
 - 5.2 Secondary personas
 - 5.3 Anti-personas
6. Product Goals & Success Criteria
 - 6.1 Product goals (12-month horizon)
 - 6.2 North Star Metric
 - 6.3 Counter-metrics (to prevent gaming)
7. Functional Requirements
 - 7.1 Authentication & Identity (REQ-AUTH)
 - 7.2 Discovery (REQ-DISC)
 - 7.3 Validation (REQ-VAL)

- 7.6 Dashboards & Workspace (REQ-DASH)
- 7.7 Integration & API (REQ-INT)
- 7.8 Billing & Plans (REQ-BIL)
- 7.9 Settings & Admin (REQ-ADMIN)
- 7.10 Onboarding & Help (REQ-OB)
- 7.11 Platform internals (REQ-PLAT)
- 8. Non-Functional Requirements
 - 8.1 Performance
 - 8.2 Availability
 - 8.3 Scalability
 - 8.4 Security & compliance
 - 8.5 Accessibility
 - 8.6 Internationalization
 - 8.7 Data retention
 - 8.8 Disaster recovery
- 9. Feature Catalogue
- 10. User Stories
 - US-FOUNDER-001 — Trend exploration
 - US-FOUNDER-002 — Pain point harvesting
 - US-FOUNDER-003 — One-page brief
 - US-FOUNDER-004 — GTM diagnosis
 - US-CORP-001 — Adjacency exploration
 - US-CORP-002 — Quarterly report
 - US-INVESTOR-001 — Deal memo
 - US-INVESTOR-002 — Thesis portfolio
 - US-CONSULT-001 — Client report
 - US-CONSULT-002 — Multi-engagement workspace
 - US-PLAT-001 — API access
 - US-PLAT-002 — Audit log
- 11. KPIs & Success Metrics
- 12. Business Model
 - 12.1 Pricing tiers
 - 12.2 Revenue model
 - 12.3 Cost structure

- 13.1 Top risks
- 13.2 Key assumptions
- 14. MVP Scope & Product Roadmap
 - 14.1 MVP definition (Q3 2027 launch)
 - 14.2 Roadmap
 - 14.3 Release criteria (per release)
- 15. Glossary & Appendix
 - 15.1 Glossary
 - 15.2 Appendix A — Document index
 - 15.3 Appendix B — Requirement → Document map
 - 15.4 Revision history

Product Requirements Document (PRD)

Document 01 — VentureMiner AI The canonical product specification. This document defines the what and why of the AI Venture Intelligence Platform. All other documents (architecture, design, engineering) derive from the requirements captured here.

Table of Contents

1. Executive Summary
2. Product Vision & Mission
3. Market Opportunity
4. Problem Statement
5. Stakeholders & User Personas
6. Product Goals & Success Criteria
7. Functional Requirements
8. Non-Functional Requirements
9. Feature Catalogue
10. User Stories
11. KPIs & Success Metrics
12. Business Model
13. Risks & Assumptions
14. MVP Scope & Product Roadmap
15. Glossary & Appendix

1. Executive Summary

VentureMiner AI is an AI Venture Intelligence Platform that turns raw, fragmented market signals into validated, monetizable SaaS opportunities.

Today's entrepreneurs, intrapreneurs, and investors spend 60–80% of their discovery time on manual research — RSS feeds, Twitter threads, Product Hunt, LinkedIn, Google Trends, ad libraries, app store charts, job boards, SEC filings. The result is a low-signal, high-bias pipeline that misses opportunities or duplicates ideas already validated by better-resourced teams.

VentureMiner AI replaces that pipeline with an autonomous, multi-agent research system that continuously:

1. Discovers candidate opportunities from public, semi-public, and licensed sources.
2. Validates each opportunity against market size, growth, demand, willingness to pay, and competitive landscape.
3. Scores opportunities on a transparent, configurable rubric (market, demand, buildability, defensibility, AI-fit).

The platform is sold as a multi-tenant SaaS to four primary personas (Section 5): independent founders, corporate innovation teams, investors, and product/strategy consultants.

1.1 One-line value proposition

From signal noise to venture-ready decisions — in hours, not quarters.

1.2 Strategic positioning

VentureMiner AI is positioned as a decision-support layer above the existing tool stack. It does not replace CRM, BI, or financial modeling tools — it provides the upstream intelligence those tools act on. Differentiation comes from three places:

- Multi-agent research depth — not one LLM prompt, but a coordinated team of specialist agents.
- Evidence discipline — every claim is cited; every score has a trace.
- Calibrated scoring — the rubric is exposed, versioned, and editable.

1.3 At a glance

Document	PRD v1.0
Total requirements	174
MVP features	22
Personas	4 primary, 2 secondary
Target launch	Q3 2027 (MVP)
Pricing tiers	4 (Free, Solo, Team, Enterprise)

2. Product Vision & Mission

2.1 Vision

A world where anyone with ambition and a credible idea can validate, refine, and launch a venture backed by the same caliber of market intelligence historically reserved for top-tier venture firms.

2.2 Mission

Build the most rigorous, transparent, and fast AI research platform for venture opportunity discovery and validation.

2.3 Guiding principles

1. Evidence over opinion. Every claim is sourced. Every score is traceable.
2. Transparent models. The rubric is a first-class artifact, not a black box.
3. Human-in-the-loop. AI accelerates; humans decide.
4. Compounding knowledge. Each run enriches the platform's corpus for the next.
5. Operator-grade reliability. If a researcher would trust it, an operator must trust it.

3. Market Opportunity

3.1 Total Addressable Market (TAM)

The global market for market intelligence and competitive intelligence software is approximately \$52B in 2026, growing at a CAGR of ~13%. The subset relevant to VentureMiner AI — opportunity discovery, market sizing, and competitive monitoring — is ~\$8.4B.

3.2 Serviceable Addressable Market (SAM)

Geographically (NA + EU + ANZ) and segmentally (SMB founders, corporate innovation, VCs, consultancies), the SAM is ~\$2.1B.

3.3 Serviceable Obtainable Market (SOM)

Realistic 5-year capture assuming the product-market fit trajectory in this PRD: \$45M ARR ≈ ~\$63M in platform spend at gross-margin-blended pricing.

3.4 Why now

- LLM capability inflection. Multi-agent orchestration and tool-use are now production-ready (2025–2026).
- Cost collapse. Inference cost per 1M tokens has dropped ~12x in 18 months.
- Source accessibility. APIs for X, Reddit, GitHub, Stripe, App Stores, SEC EDGAR, job boards, Google Trends, Similarweb, and SEMrush are mature and affordable.
- Founder saturation. Globally, ~580M people attempt a business each year; even 0.1% adoption of a \$40/mo tool is a \$28M ARR business.
- Enterprise AI budget growth. Corporate innovation teams have dedicated AI tooling budgets (avg \$310k/year) — a wedge into enterprise.

3.5 Competitive landscape

Market intelligence suites	Crayon, Klue, Kompyte	Enterprise CI, no opportunity generation
Trend monitoring	TrendWatching, Trend Hunter	Editorial, no AI
Indie research	SparkToro, Exploding Topics	Niche signals, no validation
VC deal flow	CB Insights, PitchBook	Closed-network deal sourcing
LLM research assistants	Perplexity Pro, You.com, ChatGPT Deep Research	Generic, not domain-tuned

White space: no current product offers closed-loop discovery → validation → scoring → report for SaaS opportunities, with evidence discipline and an editable rubric.

4. Problem Statement

4.1 Jobs-to-be-done

When a founder, intrapreneur, or investor sits down to identify a new opportunity, they hire a research process to do the following jobs:

- JOB-1 — Find ventures worth pursuing, faster than their peers.
- JOB-2 — Confirm the opportunity is real, growing, and monetizable.

- JOB-5 — Communicate the decision to partners, LPs, or executives.

4.2 Today, this is broken

Manual research	60–80% of discovery time	Weeks → quarters of latency
Bias	Confirmation bias, recency bias, network bias	Missed opportunities, false positives
Inconsistency	Different researchers, different rubrics	Non-comparable decisions
Opacity	"Why this score?"	Unauditable decisions
Brittleness	One spreadsheet crash = lost context	Repeated work

4.3 Why this is the moment to fix it

Until 2024, this problem resisted automation because:

- Source integration was expensive.
- LLMs hallucinated evidence.
- Orchestrating multi-step research was brittle.

By 2026, all three are solvable. VentureMiner AI is built to capitalize on this convergence.

5. Stakeholders & User Personas

5.1 Primary personas

P-1: Indie Founder (e.g. Maya)

- Demographics: Solo or two-person team, pre-seed, technical.
- Time per week on opportunity discovery: 4–10 hours.
- Pains: Idea generation, validation, GTM.
- Gains sought: Confidence, speed, differentiation.
- Buying behavior: \$30–\$80/mo self-serve; will pay annually for a deal.
- Success measure: "I shipped a v1 in 90 days from first prompt."

P-2: Corporate Innovator (e.g. Daniel)

- Demographics: Innovation scout, corporate strategy, or new-venture team inside a Fortune 1000.
- Time per week on opportunity discovery: 10–20 hours.
- Pains: Quarterly reporting, executive-ready artifacts, sourcing outside the core market.
- Gains sought: Audit trail, repeatable process, breadth.
- Buying behavior: \$25k–\$120k/year contract; security review required.

• Success measure: "I delivered a pipeline of 12 validated ediscoveries per quarter."

- Demographics: Solo capitalist or analyst at a \$50M–\$500M fund.
- Time per week on opportunity discovery: 8–15 hours.
- Pains: Sourcing off-market, thesis consistency, deal memo throughput.
- Gains sought: Defensibility of thesis, memo reuse, faster screening.
- Buying behavior: \$300–\$1,500/seat/mo.
- Success measure: "I doubled my sourced-deal volume without adding headcount."

P-4: Product / Strategy Consultant (e.g. Lin)

- Demographics: Boutique consultancy or independent, 3–30-person firms.
- Time per week on opportunity discovery: 20+ hours (per engagement).
- Pains: Replicability, client-ready reports, source provenance.
- Gains sought: Brand differentiation, deliverable speed.
- Buying behavior: \$200–\$600/seat/mo, multi-seat.
- Success measure: "I deliver 3x more research engagements per quarter."

5.2 Secondary personas

- S-1: Growth-stage founder — validating adjacencies, market expansion.
- S-2: Operator at a B2B SaaS company — competitive positioning, market monitoring.

5.3 Anti-personas

We explicitly do not build for:

- A-1: Hobbyists without willingness to act — low willingness to pay, high support load.
- A-2: Speculative trend-chasers — chase every signal, never ship.
- A-3: Pure coders who never speak to customers — research is a means, not the product.

6. Product Goals & Success Criteria

6.1 Product goals (12-month horizon)

PG-1	Reach 5,000 paid users across all tiers	CEO + Growth
PG-2	Reach \$3M ARR	CFO + Sales
PG-3	Hit a sustained NPS ≥ 50	Product + Support
PG-4	Maintain $\geq 92\%$ weekly active rate among paid users	Product
PG-5	Reduce median time-to-first-report from 6h to < 45 min	Engineering

PG-1	Reach 5,000 paid users across all tiers	CEO + Growth
PG-6	Achieve 95%+ report acceptance (operator sign-off without rewrite)	Product
PG-7	Land 25 enterprise customers (\$25k+ ACV)	Sales
PG-8	Reach SOC 2 Type II + ISO 27001	Security + Compliance
PG-9	Time-to-value $\leq$ 10 minutes from signup to first useful artifact	Product
PG-10	Achieve 80% rubric customization adoption on Team/Enterprise	Product

6.2 North Star Metric

Validated opportunities acted on per week per active user. This is the metric that ties the product to behavior we want to drive: not "reports generated" (gameable) but "opportunities a user committed to next action on."

6.3 Counter-metrics (to prevent gaming)

- Reports-per-user (cap to prevent farming)
- Report rejection rate (operator edit-after-generation)
- Subscription churn (per cohort)
- Time-to-edit (proxy for low-quality output)

7. Functional Requirements

This section enumerates functional requirements at the requirement-ID level defined in Document 00 §7.1. Each requirement has a unique ID, a one-line statement, an owner, an acceptance criterion summary, and a priority (P0 = MVP, P1 = GA, P2 = post-GA).

7.1 Authentication & Identity (REQ-AUTH)

REQ-AUTH-0001	Email/password signup with email verification	P0	AC-AUTH-0001: User can sign up, verify, and log in.
REQ-AUTH-0002	OAuth (Google, GitHub, Microsoft)	P0	AC-AUTH-0002: User can link/unlink an OAuth provider.
REQ-AUTH-0003	Magic-link sign-in	P1	AC-AUTH-0003: User can sign in via one-time link.

REQ-AUTH-0001	Email/password signup with email verification	P0	AC-AUTH-0001: User can sign up, verify, and log in.
REQ-AUTH-0004	MFA via TOTP	P1	AC-AUTH-0004: User can enroll and challenge with TOTP.
REQ-AUTH-0005	SSO (SAML 2.0 / OIDC) for Enterprise	P0 for Enterprise	AC-AUTH-0005: Admin can configure IdP; users SSO.
REQ-AUTH-0006	SCIM 2.0 provisioning	P0 for Enterprise	AC-AUTH-0006: Sync creates/updates/disables users.
REQ-AUTH-0007	Role-based access control (Owner, Admin, Member, Viewer)	P0	AC-AUTH-0007: Role gates all write actions.
REQ-AUTH-0008	Workspace + project membership	P0	AC-AUTH-0008: Users can belong to multiple workspaces.
REQ-AUTH-0009	Audit log of auth events	P0 for Enterprise	AC-AUTH-0009: Log every sign-in, role change, token issue.
REQ-AUTH-0010	Session timeout & refresh token rotation	P0	AC-AUTH-0010: 30-day refresh, 8-hour idle.

7.2 Discovery (REQ-DISC)

REQ-DISC-0001	Trend mining across configured sources	P0	AC-DISC-0001: User sees ≥ 20 weekly trends.
REQ-DISC-0002	Niche topic exploration	P0	AC-DISC-0002: User enters seed $\rightarrow$ returns ≥ 10 adjacent topics.
REQ-DISC-0003	Pain-point harvesting from public forums	P0	AC-DISC-0003: User receives cited pain quotes.
REQ-DISC-0004	Competitor-derivative opportunity generation	P0	AC-DISC-0004: User uploads competitor list $\rightarrow$ returns derivative opportunities.
REQ-DISC-0005	Source configuration (per workspace)	P1	AC-DISC-0005: Admin can enable/disable sources.

REQ-DISC-0001	Trend mining across configured sources	P0	AC-DISC-0001: User sees ≥ 20 weekly trends.
REQ-DISC-0006	Continuous background discovery ("always-on")	P1	AC-DISC-0006: New opportunities surface daily.
REQ-DISC-0007	Deduplication across sources	P0	AC-DISC-0007: Same opportunity surfaces once.
REQ-DISC-0008	Time-windowed discovery (last 7/30/90 days)	P0	AC-DISC-0008: User filters discovery window.
REQ-DISC-0009	Source attribution on every opportunity	P0	AC-DISC-0009: Each opportunity lists its sources.
REQ-DISC-0010	Saved searches	P1	AC-DISC-0010: User can save and re-run a discovery query.

7.3 Validation (REQ-VAL)

REQ-VAL-0001	Market size estimation (TAM/SAM/SOM)	P0	AC-VAL-0001: User sees triangulated estimate with confidence.
REQ-VAL-0002	Growth-rate estimation	P0	AC-VAL-0002: User sees YoY and CAGR with sources.
REQ-VAL-0003	Demand signal aggregation	P0	AC-VAL-0003: User sees search volume, social volume, and intent.
REQ-VAL-0004	Competitive landscape mapping	P0	AC-VAL-0004: User sees competitor table with positioning.
REQ-VAL-0005	Pricing benchmark collection	P0	AC-VAL-0005: User sees competitor pricing tiers.
REQ-VAL-0006	Customer review synthesis	P0	AC-VAL-0006: User sees top 5 praised and complained features.

REQ-VAL-0001	Market size estimation (TAM/SAM/SOM)	P0	AC-VAL-0001: User sees triangulated estimate with confidence.
REQ-VAL-0007	Buyer persona synthesis	P0	AC-VAL-0007: User sees 1–3 buyer personas with citations.
REQ-VAL-0008	Willingness-to-pay signal estimation	P0	AC-VAL-0008: User sees WTP range with rationale.
REQ-VAL-0009	GTM-channel diagnosis	P0	AC-VAL-0009: User sees top 3 channels used by competitors.
REQ-VAL-0010	Risk register	P0	AC-VAL-0010: User sees top 5 risks with mitigations.
REQ-VAL-0011	Validation depth setting (Quick / Standard / Deep)	P0	AC-VAL-0011: User selects depth, sees time estimate.
REQ-VAL-0012	Validation re-run with new data	P1	AC-VAL-0012: User can re-validate and diff.
REQ-VAL-0013	Human-expert override (manual evidence add)	P0	AC-VAL-0013: User can attach manual evidence to a finding.
REQ-VAL-0014	Calibration against prior decisions	P2	AC-VAL-0014: User feedback improves future estimates.

7.4 Scoring (REQ-SCORE)

REQ-SCORE-0001	Default rubric v1.0	P0	AC-SCORE-0001: Default rubric with 5 dimensions.
REQ-SCORE-0002	Custom rubric editor	P0 for Team	AC-SCORE-0002: User can add, weight, rename dimensions.
REQ-SCORE-0003	Per-dimension sub-criteria	P1	AC-SCORE-0003: User can break a dimension into sub-criteria.

REQ-SCORE-0001	Default rubric v1.0	P0	AC-SCORE-0001: Default rubric with 5 dimensions.
REQ-SCORE-0004	Rubric versioning	P0 for Team	AC-SCORE-0004: Each rubric change is versioned.
REQ-SCORE-0005	Score rationale trace	P0	AC-SCORE-0005: Every dimension score shows contributing evidence.
REQ-SCORE-0006	Score normalization across rubrics	P1	AC-SCORE-0006: Scores comparable across versions.
REQ-SCORE-0007	Score confidence intervals	P1	AC-SCORE-0007: Each score has low/mid/high estimate.
REQ-SCORE-0008	Portfolio view (multi-opportunity ranking)	P0	AC-SCORE-0008: User can sort and filter by score.
REQ-SCORE-0009	Risk-adjusted score	P1	AC-SCORE-0009: Score weighted by risk dimension.
REQ-SCORE-0010	Custom scoring weights per persona/team	P1	AC-SCORE-0010: User can save weight presets.

7.5 Reporting (REQ-RPT)

REQ-RPT-0001	One-page opportunity brief	P0	AC-RPT-0001: Generated within 60 seconds.
REQ-RPT-0002	Full validation report (10–25 pages)	P0	AC-RPT-0002: Cover, exec summary, sections, sources.
REQ-RPT-0003	Executive deck (slides)	P1	AC-RPT-0003: Exportable to PPTX/Keynote.
REQ-RPT-0004	Comparison report (≥ 2 opportunities)	P0	AC-RPT-0004: Side-by-side comparison.
REQ-RPT-0005	Report customization (chapter on/off)	P1	AC-RPT-0005: User toggles sections.

REQ-RPT-0001	One-page opportunity brief	P0	AC-RPT-0001: Generated within 60 seconds.
REQ-RPT-0006	Export to PDF, DOCX, MD, HTML	P0	AC-RPT-0006: All four formats available.
REQ-RPT-0007	Embed charts in reports	P0	AC-RPT-0007: Charts auto-rendered from data.
REQ-RPT-0008	White-label reports (Enterprise)	P0 for Enterprise	AC-RPT-0008: Logo, palette, footer configurable.
REQ-RPT-0009	Report templates	P1	AC-RPT-0009: User can save a report template.
REQ-RPT-0010	Citations and footnotes	P0	AC-RPT-0010: Every claim is footnoted.
REQ-RPT-0011	Provenance panel (sources, freshness)	P0	AC-RPT-0011: Each report lists its data freshness.
REQ-RPT-0012	Scheduled report delivery	P1	AC-RPT-0012: User can subscribe a digest.

7.6 Dashboards & Workspace (REQ-DASH)

REQ-DASH-0001	Opportunity pipeline board	P0	AC-DASH-0001: Kanban (Watching, Validating, Active, Archived).
REQ-DASH-0002	Portfolio view	P0	AC-DASH-0002: Multi-dimensional table.
REQ-DASH-0003	Saved filters	P0	AC-DASH-0003: User can save filter combos.
REQ-DASH-0004	Alerts & notifications	P0	AC-DASH-0004: User can configure alert rules.
REQ-DASH-0005	Trends dashboard	P0	AC-DASH-0005: Surface top trends per workspace.
REQ-DASH-0006	Activity feed	P0	AC-DASH-0006: Per-workspace activity timeline.

REQ-DASH-0001	Opportunity pipeline board	P0	AC-DASH-0001: Kanban (Watching, Validating, Active, Archived).
REQ-DASH-0007	Quick search (Cmd-K)	P0	AC-DASH-0007: Global search across opportunities, reports, sources.
REQ-DASH-0008	Watchlists	P1	AC-DASH-0008: User can watch keywords or competitors.
REQ-DASH-0009	Multi-workspace switching	P0	AC-DASH-0009: User can switch in one click.

7.7 Integration & API (REQ-INT)

REQ-INT-0001	REST API	P0	AC-INT-0001: OpenAPI 3.1 spec published.
REQ-INT-0002	Webhooks	P1	AC-INT-0002: Outbound event delivery.
REQ-INT-0003	Zapier integration	P2	AC-INT-0003: Public Zap.
REQ-INT-0004	Notion integration	P1	AC-INT-0004: Push opportunities and reports.
REQ-INT-0005	Slack integration	P0	AC-INT-0005: Receive alerts in Slack.
REQ-INT-0006	Linear / Jira integration	P2	AC-INT-0006: Convert opportunity into issue.
REQ-INT-0007	CSV import/export	P0	AC-INT-0007: Round-trip preserves data.
REQ-INT-0008	Google Drive / Dropbox export	P2	AC-INT-0008: Export reports to folder.
REQ-INT-0009	Public API tokens (scoped)	P0	AC-INT-0009: User can mint scoped tokens.
REQ-INT-0010	Rate-limited public API	P0	AC-INT-0010: Tier-based rate limits.

7.8 Billing & Plans (REQ-BIL)

REQ-BIL-0001	Free tier with quotas	P0	AC-BIL-0001: 5 reports/month, 1 user.
REQ-BIL-0002	Solo plan (\$39/mo)	P0	AC-BIL-0002: 50 reports/month, 1 user.
REQ-BIL-0003	Team plan (\$199/mo + \$39/seat)	P0	AC-BIL-0003: Multi-user, custom rubric, SSO-ready.
REQ-BIL-0004	Enterprise plan (custom)	P0	AC-BIL-0004: SSO, SCIM, custom legal, SLA.
REQ-BIL-0005	Stripe-backed billing	P0	AC-BIL-0005: Self-serve checkout.
REQ-BIL-0006	Annual discount (16%)	P0	AC-BIL-0006: Annual plans show savings.
REQ-BIL-0007	Usage-based overages	P1	AC-BIL-0007: Soft cap with overage.
REQ-BIL-0008	Plan upgrade/downgrade	P0	AC-BIL-0008: Pro-rated billing.
REQ-BIL-0009	Tax handling (US sales tax, EU VAT)	P0	AC-BIL-0009: Tax computed and itemized.
REQ-BIL-0010	Invoice history	P0	AC-BIL-0010: All invoices downloadable.

7.9 Settings & Admin (REQ-ADMIN)

REQ-ADMIN-0001	Workspace settings	P0	AC-ADMIN-0001: Name, logo, default rubric.
REQ-ADMIN-0002	Member management	P0	AC-ADMIN-0002: Invite, remove, role change.
REQ-ADMIN-0003	Source management	P0	AC-ADMIN-0003: Enable/disable + credentials.
REQ-ADMIN-0004	Billing & plan management	P0	AC-ADMIN-0004: Self-serve plan changes.
REQ-ADMIN-0005	Audit log (Enterprise)	P0 for Enterprise	AC-ADMIN-0005: Filterable, exportable.

REQ-ADMIN-0001	Workspace settings	P0	AC-ADMIN-0001: Name, logo, default rubric.
REQ-ADMIN-0006	API token management	P0	AC-ADMIN-0006: Mint/revoke scoped tokens.
REQ-ADMIN-0007	Custom report branding	P0 for Enterprise	AC-ADMIN-0007: Logo, palette, footer.
REQ-ADMIN-0008	Data retention configuration	P0 for Enterprise	AC-ADMIN-0008: Per-data-class retention.
REQ-ADMIN-0009	Webhook configuration	P1	AC-ADMIN-0009: Endpoint + secret.
REQ-ADMIN-0010	Compliance & DPA downloads	P0 for Enterprise	AC-ADMIN-0010: Self-serve DPA, SOC 2 report.

7.10 Onboarding & Help (REQ-OB)

REQ-OB-0001	Interactive product tour	P0	AC-OB-0001: First-run tour, skippable.
REQ-OB-0002	Sample opportunities for first run	P0	AC-OB-0002: One-click "try this".
REQ-OB-0003	Help center (in-app)	P0	AC-OB-0003: Search + articles.
REQ-OB-0004	Email support (Free, Solo)	P0	AC-OB-0004: < 48h response.
REQ-OB-0005	In-app chat (Team)	P0	AC-OB-0005: < 8h business hours.
REQ-OB-0006	Dedicated CSM (Enterprise)	P0 for Enterprise	AC-OB-0006: Named CSM, monthly QBR.
REQ-OB-0007	In-app feedback widget	P0	AC-OB-0007: One-click feedback.
REQ-OB-0008	Status page	P0	AC-OB-0008: Public status.history.
REQ-OB-0009	Changelog & release notes	P0	AC-OB-0009: In-app + RSS.
REQ-OB-0010	Onboarding checklist	P0	AC-OB-0010: Tracks first 5 actions.

7.11 Platform internals (REQ-PLAT)

REQ-PLAT-0001	Multi-tenant isolation	P0	AC-PLAT-0001: Row-level isolation per workspace.
REQ-PLAT-0002	Configurable data residency	P0 for Enterprise	AC-PLAT-0002: US, EU, APAC.
REQ-PLAT-0003	Encryption at rest (AES-256)	P0	AC-PLAT-0003: All persistent storage.
REQ-PLAT-0004	Encryption in transit (TLS 1.3)	P0	AC-PLAT-0004: All network traffic.
REQ-PLAT-0005	Key management (KMS)	P0	AC-PLAT-0005: Per-tenant DEKs.
REQ-PLAT-0006	Background job orchestration	P0	AC-PLAT-0006: Async pipelines survive restarts.
REQ-PLAT-0007	Idempotency keys on write APIs	P0	AC-PLAT-0007: Duplicate retries are no-ops.
REQ-PLAT-0008	Event sourcing for AI pipeline	P0	AC-PLAT-0008: Every AI step is replayable.
REQ-PLAT-0009	Rate limiting per workspace	P0	AC-PLAT-0009: Tier-based limits.
REQ-PLAT-0010	Health endpoints	P0	AC-PLAT-0010: /healthz, /readyz.

8. Non-Functional Requirements

8.1 Performance

NFR-PERF-001	Time to first dashboard render	< 1.5s p75
NFR-PERF-002	One-page brief generation	< 60s p75
NFR-PERF-003	Full report generation (Standard depth)	< 8 min p75
NFR-PERF-004	Search response	< 300ms p75
NFR-PERF-005	API p95 latency	< 500ms (sync) / async otherwise

8.2 Availability

- Production SLO: 99.9% monthly availability for paid tiers.
- Free tier SLO: 99.5%.

8.3 Scalability

- Support 10,000 concurrent users at GA, 50,000 at Year 2.
- 1M opportunities per workspace maximum.
- 100M documents in the platform knowledge base.

8.4 Security & compliance

- SOC 2 Type II at GA; ISO 27001 by Month 18.
- GDPR + CCPA compliant at launch.
- HIPAA not in scope for v1.
- All customer data encrypted at rest and in transit.
- Penetration test annually + on major releases.

8.5 Accessibility

- WCAG 2.1 AA across all user-facing pages.
- Keyboard navigability for every interaction.
- Screen-reader compatibility (NVDA, VoiceOver).
- Color contrast $\geq 4.5:1$ for text.

8.6 Internationalization

- v1 supports English (en-US).
- v2 adds Spanish, German, Japanese.

8.7 Data retention

- User content retained while account is active.
- Soft-delete with 30-day recovery window.
- Hard delete after 90 days, fully purged.
- Configurable per-tenant for Enterprise.

8.8 Disaster recovery

- RPO: 1 hour.
- RTO: 4 hours.
- Multi-region backups (US-EAST primary, US-WEST standby).

9. Feature Catalogue

The following 35 features (FEAT-001..FEAT-035) compose the MVP + GA product. Each is mapped to the requirements it satisfies.

FEAT-001	Email/OAuth sign-up	REQ-AUTH-0001..0002	✓	
FEAT-002	SSO + SCIM (Enterprise)	REQ-AUTH-0005..0006	✓	
FEAT-003	MFA	REQ-AUTH-0004		✓
FEAT-004	RBAC	REQ-AUTH-0007	✓	
FEAT-005	Multi-workspace	REQ-AUTH-0008	✓	
FEAT-006	Trend mining	REQ-DISC-0001	✓	
FEAT-007	Niche exploration	REQ-DISC-0002	✓	
FEAT-008	Pain harvesting	REQ-DISC-0003	✓	
FEAT-009	Competitor-derivative	REQ-DISC-0004	✓	
FEAT-010	Continuous discovery	REQ-DISC-0006		✓
FEAT-011	Saved searches	REQ-DISC-0010		✓
FEAT-012	Market size	REQ-VAL-0001	✓	
FEAT-013	Growth rate	REQ-VAL-0002	✓	
FEAT-014	Demand signals	REQ-VAL-0003	✓	
FEAT-015	Competitive map	REQ-VAL-0004	✓	
FEAT-016	Pricing benchmarks	REQ-VAL-0005	✓	
FEAT-017	Review synthesis	REQ-VAL-0006	✓	
FEAT-018	Buyer personas	REQ-VAL-0007	✓	
FEAT-019	WTP estimation	REQ-VAL-0008	✓	
FEAT-020	GTM diagnosis	REQ-VAL-0009	✓	
FEAT-021	Risk register	REQ-VAL-0010	✓	
FEAT-022	Validation depth	REQ-VAL-0011	✓	
FEAT-023	Default rubric	REQ-SCORE-0001	✓	
FEAT-024	Custom rubric	REQ-SCORE-0002	✓	
FEAT-025	Rubric versioning	REQ-SCORE-0004		✓
FEAT-026	Score rationale	REQ-SCORE-0005	✓	
FEAT-027	Portfolio view	REQ-SCORE-0008	✓	
FEAT-028	One-page brief	REQ-RPT-0001	✓	
FEAT-029	Full report	REQ-RPT-0002	✓	
FEAT-030	Comparison report	REQ-RPT-0004	✓	
FEAT-031	White-label	REQ-RPT-0008	✓ (Ent)	

FEAT-001	Email/OAuth sign-up	REQ-AUTH-0001..0002	✓	
FEAT-032	Pipeline board	REQ-DASH-0001	✓	
FEAT-033	Alerts	REQ-DASH-0004	✓	
FEAT-034	Slack integration	REQ-INT-0005	✓	
FEAT-035	Public REST API	REQ-INT-0001		✓

MVP scope: 22 features (✓ in MVP column). The remaining 13 are post-MVP / GA-only.

10. User Stories

A sample of 12 representative user stories. The full backlog lives in the project's tracking tool; these are the spec-defining ones.

US-FOUNDER-001 — Trend exploration

As Maya (indie founder), I want to enter a broad interest like "AI for solopreneurs" and see the top 20 trending sub-niches so that I can pick one to dig into.

Acceptance Criteria:

- AC-US-FOUNDER-001-1: System returns ≥ 20 sub-niches ranked by composite trend score.
- AC-US-FOUNDER-001-2: Each sub-niche lists 3+ supporting sources.
- AC-US-FOUNDER-001-3: User can pin a sub-niche to start validation.

US-FOUNDER-002 — Pain point harvesting

As Maya, I want to see the top complaints and unmet needs in my chosen niche from public forums so that I can ground my product idea in real demand.

Acceptance Criteria:

- AC-US-FOUNDER-002-1: System aggregates from ≥ 3 configured sources (Reddit, HN, X, G2).
- AC-US-FOUNDER-002-2: Each pain point has ≥ 1 direct quote and a permalink.
- AC-US-FOUNDER-002-3: User can save pains to a "Discovery notes" panel.

US-FOUNDER-003 — One-page brief

As Maya, I want a one-page brief on a sub-niche I selected so I can share it with a potential co-founder in 60 seconds.

Acceptance Criteria:

- AC-US-FOUNDER-003-1: Brief includes: opportunity, target user, market size band, top 3 competitors, score, top 3 risks, citations.
- AC-US-FOUNDER-003-2: Generation completes in $< 60s$ p75.
- AC-US-FOUNDER-003-3: User can export to PDF or share a read-only link.

US-FOUNDER-004 — GTM diagnosis

As Maya, I want to know which channels competitors use to acquire users so I can pick a launch channel.

Acceptance Criteria:

- AC-US-FOUNDER-004-1: GTM panel lists top 3 channels with evidence.
- AC-US-FOUNDER-004-2: Each channel has an estimated CAC band.

US-CORP-001 — Adjacency exploration

As Daniel (corporate innovator), I want to explore adjacencies to our core business so we can identify new product lines.

Acceptance Criteria:

- AC-US-CORP-001-1: User can enter a base business descriptor; system returns ≥ 10 adjacencies.
- AC-US-CORP-001-2: Each adjacency is mapped to existing capabilities and customer segments.

US-CORP-002 — Quarterly report

As Daniel, I want to deliver a quarterly board report summarizing 10 validated opportunities so leadership can decide on resourcing.

Acceptance Criteria:

- AC-US-CORP-002-1: Report is a multi-section PDF with cover, exec summary, ranked opportunities, top 3 recommendations.
- AC-US-CORP-002-2: White-label branding is applied.
- AC-US-CORP-002-3: Report generation completes in < 30 minutes for 10 opportunities.

US-INVESTOR-001 — Deal memo

As Aria (investor), I want a deal-memo-style brief on a startup that surfaced in my feed so I can decide whether to take a meeting.

Acceptance Criteria:

- AC-US-INVESTOR-001-1: Memo includes: thesis fit, market context, comparable rounds, founder-style questions.
- AC-US-INVESTOR-001-2: User can attach the memo to a CRM record via API.

US-INVESTOR-002 — Thesis portfolio

As Aria, I want to see my entire pipeline scored against my thesis so I can balance the portfolio.

Acceptance Criteria:

- AC-US-INVESTOR-002-1: System scores each opportunity against saved thesis criteria.
- AC-US-INVESTOR-002-2: User can see concentration / coverage heatmap.

US-CONSULT-001 — Client report

As Lin (consultant), I want a branded client report so I can deliver it under my firm's identity.

Acceptance Criteria:

US-CONSULT-002 — Multi-engagement workspace

As Lin, I want one workspace per client so that work does not leak between engagements.

Acceptance Criteria:

- AC-US-CONSULT-002-1: User can create up to 5 workspaces (Free), 25 (Team), unlimited (Enterprise).
- AC-US-CONSULT-002-2: Members can be added to a specific workspace only.

US-PLAT-001 — API access

As any user, I want a token-based API so I can integrate the platform with my workflow.

Acceptance Criteria:

- AC-US-PLAT-001-1: User can mint a token scoped to specific resources.
- AC-US-PLAT-001-2: Token can be revoked.

US-PLAT-002 — Audit log

As a security officer, I want to see who did what so I can answer audit questions.

Acceptance Criteria:

- AC-US-PLAT-002-1: Audit log captures sign-ins, role changes, exports, deletes.
- AC-US-PLAT-002-2: Audit log is exportable to CSV.

11. KPIs & Success Metrics

MRR	\$250k → \$3M	12 months	Stripe
Paid users	5,000	12 months	App
Activation rate	70% (signup → first report)	30 days	App
Weekly active rate (paid)	92%	Continuous	App
Report acceptance rate	95%	Continuous	App
Time-to-first-report	45 min median	Continuous	App
NPS	50+	Quarterly	Survey
Logo churn (monthly)	< 3%	Continuous	Stripe
Net revenue retention	115%+	Annual	Stripe
Source coverage	95% of priority sources live	Quarterly	App

12.1 Pricing tiers

Free	\$0	1	5	No	No	None
Solo	\$39/mo	1	50	No	No	Standard
Team	\$199 + \$39/seat	5 included	300	Yes	Add-on	Standard
Enterprise	Custom	Unlimited	Custom	Yes	Yes	99.9%

12.2 Revenue model

- Subscription is 88% of revenue.
- Usage overages are 7% (driven by Team+ power users).
- Enterprise services (custom integrations, training) are 5%.

12.3 Cost structure

- Inference (LLM + tools) is the largest variable cost; targeted at < 30% of revenue.
- Engineering & product dominate fixed cost.
- Sales & marketing target CAC payback < 14 months.

12.4 Funding posture

- Seed-to-Series-A transition targeted at \$4M ARR.
- Series B at \$12M ARR for enterprise GTM build-out.

13. Risks & Assumptions

13.1 Top risks

RISK-001	LLM hallucination in evidence	High	High	Multi-agent verification + citation requirements + operator sign-off
RISK-002	Source API cost escalation	Medium	High	Multi-source fallback + cache + cost dashboards
RISK-003	Enterprise sales cycle > 6 months	Medium	Medium	Land-and-expand with Team tier as wedge
RISK-004	Regulatory change (AI Act, etc.)	Medium	Medium	Continuous compliance review + data-residency

RISK-001	LLM hallucination in evidence	High	High	Multi-agent verification + citation requirements + operator sign-off
RISK-005	Big tech competitor enters	High	High	Speed-to-feature + multi-agent depth + community
RISK-006	IP / copyright on training data	Medium	High	Source provenance + license audit + indemnity
RISK-007	Founding team concentration	Medium	High	Hire aggressively + documentation culture
RISK-008	Runway burn	Low	High	Conservative seed + revenue focus

13.2 Key assumptions

- ASP-001: Inference cost continues to fall $\geq 20\%$ YoY.
- ASP-002: Founders will pay \$39+/mo for a tool that materially reduces research time.
- ASP-003: Enterprise security reviews can be passed within 60 days of contracting.
- ASP-004: A 4-person founding team can ship MVP in 9 months.

14. MVP Scope & Product Roadmap

14.1 MVP definition (Q3 2027 launch)

- 22 features (Section 9).
- 4 personas supported at basic depth.
- Web app + REST API.
- English only.
- US/EU data residency.
- Stripe billing, 4 tiers.

14.2 Roadmap

Q3 2026	PRD locked, architecture baseline, team complete (8 FTE)
Q4 2026	Alpha: discovery + basic validation; internal dogfood

Q3 2026	PRD locked, architecture baseline, team complete (8 FTE)
Q1 2027	Beta: full MVP features behind a flag; 50 design partners
Q2 2027	Hardening: security review, pricing validation, content
Q3 2027	Public launch (Free + Solo + Team); ProductHunt, HN
Q4 2027	Enterprise plan launches; SOC 2 Type I
Q1 2028	i18n (es, de, ja); continuous discovery; webhook v1
Q2 2028	SOC 2 Type II; Notion + Linear; cohort analytics
Q3 2028	API v2; multi-tenant rubric library; 25 enterprise customers
Q4 2028	ISO 27001; 5,000 paid users; \$3M ARR target

14.3 Release criteria (per release)

- All acceptance criteria for the release scope pass.
- Performance NFRs (Section 8.1) green for 7 consecutive days.
- No open P0/P1 security issues.
- Documentation updated and indexed.
- Customer comms drafted (changelog, email, in-app banner).

15. Glossary & Appendix

15.1 Glossary

Validation	Confirming an opportunity is real, growing, and monetizable.
Discovery	Generating candidate opportunities from raw signals.
Scoring	Computing a numeric rating from a configurable rubric.
Rubric	The set of dimensions, weights, and criteria used to score.
Adjacency	A business area near (but not identical to) an existing one.
WTP	Willingness to pay.
GTM	Go-to-market.

Validation	Confirming an opportunity is real, growing, and monetizable.
TAM / SAM / SOM	Total / Serviceable / Serviceable-Obtainable market.
North Star	The single user-behavior metric that defines product success.

15.2 Appendix A — Document index

00	Documentation Governance	docs/00-Governance/00_documentation_governance.md
01	PRD (this document)	docs/01-PRD/01_product_requirements_document.md
02	TRD	docs/02-TRD/02_technical_requirements_document.md
03	Application Flow	docs/03-ApplicationFlow/03_application_flow.md
04	UI/UX Spec	docs/04-UIUX/04_uiux_design_specification.md
05	Backend & DB	docs/05-Backend/05_backend_database_architecture.md
06	Roadmap	docs/06-Roadmap/06_implementation_roadmap.md
07..18	AI Architecture	docs/07-AI-Architecture/
19..30	Engineering	docs/08-Engineering/

15.3 Appendix B — Requirement → Document map

AUTH	10	PRD (here), TRD, Security
DISC	10	PRD, AI Architecture, Multi-Agent, Plugin
VAL	14	PRD, Research Pipeline, Scoring
SCORE	10	PRD, Scoring Engine
RPT	12	PRD, Report Generation
DASH	9	PRD, UI/UX, Application Flow
INT	10	PRD, API Specs
BIL	10	PRD, Backend
ADMIN	10	PRD, Backend, Security
OB	10	PRD, UI/UX
PLAT	10	PRD, Backend, DevOps

15.4 Revision history

v0.1	2026-07-20	Doc Team	Initial outline
v0.5	2026-07-20	Doc Team	All 15 chapters drafted
v1.0	2026-07-20	Doc Team	First approved version

End of Document 01 — PRD. This document is the upstream source of truth for every requirement ID, persona, and KPI used elsewhere in the suite.