

VentureMiner AI

AI Operations & Evaluation

Document: 07-AI-Architecture / 18_ai_operations_evaluation

Version: v1.0

Date: 2026-07-20

Author: VentureMiner AI Documentation Team

Status: Approved

Project: VentureMiner AI — AI Venture Intelligence Platform

Table of Contents

Document 18 — AI Operations & Evaluation

Table of Contents

1. Purpose & Scope
2. AI observability
 - 2.1 Spans
 - 2.2 Metrics
 - 2.3 Logs
3. Evaluation framework
 - 3.1 Three layers
 - 3.2 Golden test sets
 - 3.3 Online metrics
4. Calibration
5. Drift detection
6. Incident response
7. Cost management
8. Continuous improvement
9. Safety & policy
10. Appendix
 - 10.1 Revision history
 - 10.2 Cross-references

Document 18 — AI Operations & Evaluation

How we operate the AI plane in production. This document defines the day-2 concerns: monitoring, evaluation, calibration, incident response, and continuous improvement.

Table of Contents

1. Purpose & Scope
2. AI observability
3. Evaluation framework
4. Calibration
5. Drift detection
6. Incident response
7. Cost management
8. Continuous improvement
9. Safety & policy
10. Appendix

1. Purpose & Scope

This document is the operational contract for the AI plane. It defines what we monitor, how we evaluate, how we respond to incidents, and how we improve the system over time.

2. AI observability

2.1 Spans

Every AI plane action is a span:

- `agent.run` — a complete run.
- `agent.step` — a single node execution.
- `rag.retrieve` — a retrieval call.
- `mcp.call` — a tool call.
- `llm.invoke` — a model call.
- `verifier.check` — a verification pass.

2.2 Metrics

- Per agent: call count, error rate, p50/p95/p99 latency, cost.
- Per tool: call count, error rate, p95 latency, cost, schema-violation rate.
- Per LLM: token in/out, cost, finish reason, fallback rate.
- Per run: total cost, total latency, verifier pass rate, retry count.

2.3 Logs

- LLM prompts and outputs logged at debug (sampled 100% for AI plane).
- Tool inputs/outputs logged at debug (PII-redacted).

3. Evaluation framework

3.1 Three layers

- Offline: golden test sets, run weekly.
- Shadow: new model/prompt runs in parallel; compared against current.
- Online: A/B tests on user-facing surfaces.

3.2 Golden test sets

- 500 labeled opportunities for validation.
- 200 labeled opportunities for scoring.
- 50 sample reports reviewed by experts.
- 100 adversarial prompts (hallucination, policy, PII).

3.3 Online metrics

- Report acceptance rate.
- Score re-edit rate.
- Re-run rate.
- Time-to-first-useful-artifact.
- Tool failure rate per user.

4. Calibration

- Monthly: score calibration against expert set (Document 16 §9).
- Quarterly: full rubric review.
- Continuous: user feedback loop feeds a calibration dataset (with consent).

5. Drift detection

- Input drift: changes in source data shape (e.g. X API breaking).
- Output drift: verifier pass rate falling, or model output distribution shifting.
- User drift: acceptance rate falling over a 7-day window.

Detection:

- Statistical tests (KS test) on embedding distributions.
- Threshold alerts on metrics.

- Roll back to last known-good model.
- Trigger rubric or prompt review.
- Notify the on-call AI lead.

6. Incident response

SEV-1	AI plane down or producing unsafe output	Page on-call; freeze deploys; revert
SEV-2	Significant quality drop or cost spike	Notify AI lead; investigate within 2h
SEV-3	Single agent degraded	Investigate within 24h
SEV-4	Cosmetic / minor	Backlog

Runbooks live in Document 28 (Operations).

7. Cost management

- Per-workspace budgets with soft caps and overage alerts.
- Per-run budgets enforced by the orchestrator.
- Cost attribution to model, tool, run, and user.
- Daily dashboards of cost per active user.
- Cost anomaly detection — alerts on $> 2\sigma$ deviation.

8. Continuous improvement

- Weekly: review online metrics.
- Bi-weekly: review agent outputs from production.
- Monthly: rubric and prompt review.
- Quarterly: full eval-set refresh and major prompt overhaul.
- Yearly: strategy review — new models, new techniques.

9. Safety & policy

- PII detection at AI plane boundary; redaction in prompts and logs.
- Policy engine filters disallowed content.
- Adversarial testing monthly.
- Red-team quarterly.
- Audit log of every policy decision.

10. Appendix

10.1 Revision history

v0.5	2026-07-20	Doc Team	All sections drafted
v1.0	2026-07-20	Doc Team	First approved version

10.2 Cross-references

- Multi-Agent: Document 08.
- RAG: Document 10.
- MCP: Document 12.
- Scoring Engine: Document 16.
- Security: Document 21.
- Operations Guide: Document 28.

End of Document 18 — AI Operations & Evaluation. The AI plane is operated, not deployed and forgotten.